

CMPT 489 / 980 Program Synthesis

Final Project

Phase I is due by 11:59pm PT on Wednesday Nov 8, 2023. Phase II is due by 11:59pm PT on Tuesday Dec 5, 2023. Please submit them to Canvas on time. No late submission is accepted.

Requirements:

- This project must be your own work. No collaboration is permitted.
- The programming language of this project is Java 11.
- You can learn the code on slides and start from it.
- You can use third-party libraries but not existing synthesizers. However, you can implement the algorithms in existing synthesizers by yourself.

1 Problem Description

Consider the following context-free grammar G

$$\begin{aligned} E &::= \text{Ite}(B, E, E) \mid \text{Add}(E, E) \mid \text{Multiply}(E, E) \mid x \mid y \mid z \mid 1 \mid 2 \mid 3 \\ B &::= \text{Lt}(E, E) \mid \text{Eq}(E, E) \mid \text{And}(B, B) \mid \text{Or}(B, B) \mid \text{Not}(B) \\ x, y, z &\in \mathbf{Variables} \quad 1, 2, 3 \in \mathbf{Constants} \end{aligned}$$

Here, E is the start symbol. E and B are non-terminals; all other symbols are terminals. The meaning of terminal symbols are self-explanatory. Specifically, **Ite** is the if-then-else operator. **Add** is the addition (+) operator. **Multiply** is the multiplication (*) operator. x, y, z are integer variables. $1, 2, 3$ are integer constants. **Lt** is the less-than (<) operator. **Eq** is the equals (==) operator. **And** is the logical conjunction (&&). **Or** is the logical disjunction (||). **Not** is the logical negation (!).

In this project, you need to write an example-based program synthesizer in Java. Specifically, the synthesizer takes as input a list of input-output examples and the context-free grammar G and produces as output an implementation of $f(x, y, z)$ in the language of G such that $f(x, y, z)$ is consistent with the provided examples. You can assume f only uses three variables x, y, z , and all their types are **Int**. The return type of f is also **Int**. If the synthesis succeeds, your program should print the program, e.g., `Add(Add(y, z), x)`, to the console. Otherwise, if the synthesis fails, the program should print `null`.

2 Codebase

A codebase is provided as the starting point. It contains the basic framework for the synthesizer. Details are explained as follows.

Package `synth.cfg`

This package defines the data structure for the context-free grammar (CFG). It has the following classes

- **Symbol**. An abstract class for symbols in the CFG.
- **Terminal**. A subclass of **Symbol** that corresponds to terminals in the CFG.
- **NonTerminal**. A subclass of **Symbol** that corresponds to non-terminals in the CFG.

- **Production.** A class for productions in the CFG. A production is of the form

$$\text{ReturnSymbol} ::= \text{Operator}(\text{ArgSymbol}, \dots, \text{ArgSymbol})$$

- **CFG.** A class for representing the CFG. The most important method is `getProductions`, which takes as input a non-terminal symbol N and returns as output a list of all productions with N being the left-hand-side of the production.

Package `synth.core`

This package contains the classes for synthesizers, examples, programs, and interpreters.

- **ASTNode.** A class for general Abstract Syntax Tree (AST) nodes. The `symbol` field corresponds to the symbol in the CFG. The `children` field corresponds to the children nodes.
- **Program.** A class for representing a program. It only has one field `root`, which is the root node of the corresponding AST.
- **Example.** A class that defines the data structure of an example. The `input` field is a map from variable names to their values. The `output` field is the output value.
- **Interpreter.** A class that defines an interpreter of the language of G . The most important method is the static method `evaluate`, which takes as input a program and an environment and returns as output the evaluation result. The environment is essentially a map from variable names to their values, just like the `input` field of `Example`. Concrete examples on how to use `Interpreter.evaluate` are provided in the test class `synth.core.InterpreterTests`.
- **ISynthesizer.** An interface that defines the input and output of a synthesizer. The inputs are a CFG and a list of examples. The output is a program.
- **TopDownEnumSynthesizer.** A top-down enumerative synthesizer that implements the `ISynthesizer` interface. **You need to implement this class.**

Package `synth.util`

This package contains the utility classes and methods.

- **FileUtils.** A class for file operations. The `readLinesFromFile` static method reads a file into a list of strings, where each line of the file is a string.
- **Parser.** A class for parsing the examples. The `parseAnExample` static method parses text of the form “`x=a, y=b, z=c -> d`” to an object of class `Example`. The `parseAllExamples` static method parses a list of examples from a list of strings, where each string corresponds to an example. It ignores empty strings.

Class `synth.Main`

The main class of the framework. It has two methods.

- **main.** It is the entry of the program. It takes one command-line argument `args[0]` for the path to the examples file.
- **buildCFG.** It builds the CFG G in Section 1.

Tests

JUnit tests are provided in the `test` directory. You are welcome to add more!

- **`synth.core.InterpreterTests`.** It contains several unit tests for the interpreter, which is also helpful for understanding the usage of the interpreter.

Other Files

- `pom.xml`. The configuration file for Maven.
- `examples.txt`. A sample examples file.

3 Compilation and Execution

Compilation. This codebase uses the Maven build system. Suppose you enter the `Synth` directory, the project can be easily compiled with one command

```
$ mvn package
```

Then you should be able to see the message “BUILD SUCCESS”. A directory called `target` will be created and a jar file called `synth-1.0.jar` will be generated inside the `target`.

Execution. In the `Synth` directory, you can execute the program using command

```
$ java -cp lib:target/synth-1.0.jar synth.Main <path-to-examples-file>
```

where `<path-to-examples-file>` is the path to the examples file. For example, you can run

```
$ java -cp lib:target/synth-1.0.jar synth.Main examples.txt
```

You will see a runtime exception with message “To be implemented”, because the synthesizer is not implemented yet. After you finish implementing the synthesizer, you should see something like (not unique)

```
Add(Add(y, z), x)
```

4 Phase I

In Phase I, you need to implement a top-down enumerative synthesizer in `synth.core.TopDownEnumSynthesizer`.

Deliverable

A zip file called `Phase1.Firstname.Lastname.zip` that contains at least the followings:

- The entire `Synth` directory. You can change existing code if you want, but please make sure the project can be compiled and executed as explained in Section 3.
- A short report (1-2 pages) called `Phase1.Firstname.Lastname.pdf` that explains the design choices, features, tests, issues (if any), and anything else that you want to explain about your program.

5 Phase II

In Phase II, you can implement any synthesis algorithm that improves the performance of the synthesizer on the same problem. **You also need to create a small benchmark set and evaluate your algorithm over the benchmarks.**

A zip file called `Phase2.Firstname.Lastname.zip` that contains at least the followings:

- The entire `Synth` directory. You can change existing code if you want, but please make sure the project can be compiled and executed as explained in Section 3.
- A long report (5-6 pages) called `Phase2.Firstname.Lastname.pdf` that explains the **algorithms**, **benchmarks**, **evaluation results**, design choices, features, tests, issues (if any), and anything else that you want to explain about your program.